

```
In [ ]: #Is It Bread? Grain Identifier
#By: Alicia Jin, Alice Liu, Jemima Chong & Bethel Yared

#Training a convolutional neural network to classify images of bread into specific catagories
```

Colab Link: [https://colab.research.google.com/drive/1pkFJD7-7jm\\_adH-XucnpVEmR2OimRc\\_x?usp=sharing](https://colab.research.google.com/drive/1pkFJD7-7jm_adH-XucnpVEmR2OimRc_x?usp=sharing)

```
In [ ]: #importing libraries
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

#useful libraries - extra options
import time
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
from torchvision import datasets, models, transforms
import time
import os
import matplotlib.pyplot as plt # Plotting graphs
from torch.utils.data import TensorDataset
```

```
In [ ]: #Mounting Google Drive as many files will be generated
from google.colab import drive
drive.mount('/content/gdrive', force_remount=True)
```

Mounted at /content/gdrive

```
In [ ]: #for splitting folders into train, validation and test set
! pip install split-folders
import splitfolders
```

Collecting split-folders  
 Downloading split\_folders-0.5.1-py3-none-any.whl (8.4 kB)  
 Installing collected packages: split-folders  
 Successfully installed split-folders-0.5.1

```
In [ ]: #Defining the path to the dataset
data_directory = '/content/gdrive/MyDrive/APS360 Project Group/Data/Images'
#test data will be 80% of all data, validaion and test 10% each
splitfolders.ratio(data_directory, 'all_data_splitted', seed=1337, ratio = (0.8, 0.1, 0.1))
```

Copying files: 1883 files [01:02, 30.35 files/s]

```
In [ ]: new_path = '/content/all_data_splitted'
train_dir = os.path.join(new_path, 'train/')
test_dir = os.path.join(new_path, 'test/')
val_dir = os.path.join(new_path, 'val/')

#Transform the images
data_transform = transforms.Compose([
    transforms.Resize((224, 224)), # Resize the images to a consistent size
    transforms.ToTensor(), # Convert images to PyTorch tensors
])

train_data = datasets.ImageFolder(train_dir, transform=data_transform)
val_data = datasets.ImageFolder(val_dir, transform=data_transform)
test_data = datasets.ImageFolder(test_dir, transform=data_transform)

#Classes: A list of strings denoting the bread depicted in each image
classes = ('Bagel', 'Baguette', 'Brioche', 'Challah', 'Cheese Bread', 'Ciabatta', 'Cinnamon Raisin', 'Focaccia', 'Wheat Grain Bread', 'White Bread')

# define dataloader parameters
batch_size = 256
num_workers = 0

# prepare data loaders
train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                           num_workers=num_workers, shuffle=True)
val_loader = torch.utils.data.DataLoader(val_data, batch_size=batch_size,
                                         num_workers=num_workers, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size,
                                           num_workers=num_workers, shuffle=True)
```

```
In [ ]: #for reference, total number of images in each set
print(len(train_data))
print(len(test_data))
print(len(val_data))
```

1500  
 196  
 184

```
In [ ]: #Display some data examples
dataiter = iter(train_loader)
images, labels = next(dataiter)
images = images.numpy()

fig = plt.figure(figsize=(25, 4))
for idx in np.arange(10):
    ax = fig.add_subplot(2, int(10/2), idx+1, xticks=[], yticks=[])
    plt.imshow(np.transpose(images[idx], (1, 2, 0)))
    ax.set_title(classes[labels[idx]])
```

Cinnamon Raisin



Ciabatta



Cinnamon Raisin



Focaccia



Cinnamon Raisin



Bagel



Cinnamon Raisin



Cheese Bread



Focaccia



Ciabatta



## Baseline Model

The Support Vector Machine (SVM) will be used to validate the performance of the primary model. Method is referenced from <https://medium.com/analytics-vidhya/image-classification-using-machine-learning-support-vector-machine-svm-dc7a0ec92e01>, <https://scikit-learn.org/stable/modules/generated/sklearn.metrics.ConfusionMatrixDisplay.html>, and tutorial 2.

```
In [ ]: from sys import float_info
import pandas as pd
import os
from skimage.io import imread
import numpy as np
import matplotlib.pyplot as plt
from skimage.transform import rescale, resize, downscale_local_mean

from PIL import Image
Image.MAX_IMAGE_PIXELS = None
import cv2

flat_data_arr = []
target_arr = []

classes = ('Bagel', 'Baguette', 'Brioche', 'Challah', 'Cheese Bread', 'Ciabatta', 'Cinnamon Raisin', 'Focaccia', 'Wheat Grain Bread', 'White Bread')

#load all data

data_directory = '/content/gdrive/MyDrive/APS360 Project Group/Data/Images'

for n in classes:
    print(f'loading... class : {n}')
    path = os.path.join(data_directory, n)
    for img in os.listdir(path):
        img_array = np.array(cv2.imread(os.path.join(path, img))) #load image data into img_array
        img_array = img_array.astype('float64')
        img_resized = resize(img_array, (20,20,3)) #resize all images to 25*25 in RGB
        flat_data_arr.append(img_resized.flatten()) #add flattened images to flat_data_arr
        target_arr.append(classes.index(n)) #add classes as numbers to target_arr
    print(f'loaded class:{n} succesfully')

In [ ]: #convert flat_data_arr and target_arr to numpy arrays
flat_data = np.array(flat_data_arr)
target = np.array(target_arr)

#define dataframe
df = pd.DataFrame(flat_data)
df = df.fillna(0) #replace NaN values with 0
df['Target'] = target
x=df.iloc[:, :-1] #input data
y=df.iloc[:, -1] #output data

In [ ]: #accelerate training
! pip install scikit-learn-intelex
from sklearnex import patch_sklearn
patch_sklearn()
from sklearn.svm import SVC

In [ ]: #define svm model
from sklearn import svm
from sklearn.model_selection import GridSearchCV
param_grid={'C':[0.1,1,10,100], 'gamma':[0.0001,0.001,0.1,1], 'kernel':['rbf', 'poly']} #either rbf or poly kernel will be used for better accuracy
svc=svm.SVC(probability=True)
model=GridSearchCV(svc,param_grid, error_score='raise') #gridsearchcv chooses the best performing parameters defined in param_grid

In [ ]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y,train_size = 0.8, shuffle = True) #split to 80% and 20%
x_val, x_test2, y_val, y_test2 = train_test_split(x_test, y_test, test_size=0.5, shuffle = True) #further split 20% into val and test

# normalize data between 0 and 1
from sklearn import preprocessing
norm_x_train = preprocessing.normalize(x_train)
norm_x_val = preprocessing.normalize(x_val)
norm_x_test2 = preprocessing.normalize(x_test2)

In [ ]: # train in batches to reduce memory
batch_size = 251
#batch size should be divisible by length of train data for better accuracy, which ensures all classes are in each batch

#training loop to iterate through batches
for i in range(0, len(norm_x_train), batch_size):
    x_train_batch = norm_x_train[i:i+batch_size]
    y_train_batch = y_train[i:i+batch_size]
    #train model with training data
    model.fit(x_train_batch, y_train_batch)
```


```
In [ ]: #validate and test model
y_pred_val = model.predict(norm_x_val) #prediction with model using validation data
y_pred_test = model.predict(norm_x_test2) #prediction with model using test data

from sklearn import metrics
from sklearn.metrics import accuracy_score
#obtain accuracy
val_accuracy = accuracy_score(y_val, y_pred_val) #compare accuracy of real vs predicted labels
test_accuracy = accuracy_score(y_test2, y_pred_test)
print(f'The validation accuracy for this model is {val_accuracy*100}%')
print(f'The test accuracy for this model is {test_accuracy*100}%')

#plot confusion matrix for testing data
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay


cm = confusion_matrix(y_test2, y_pred_test)
disp = ConfusionMatrixDisplay(confusion_matrix = cm)
disp.plot()
plt.show()
```

The validation accuracy for this model is 23.93617021276596%  
The test accuracy for this model is 25.396825396825395%



```
In [ ]: #qualitative results
#choose one random image and obtain file path to be predicted

test_image_path = '/content/gdrive/MyDrive/APS360 Project Group/Data/Images/Bagel/plain+bagel.jpg'
img=imread(test_image_path)
plt.imshow(img)
plt.show()
img_resize=resize(img,(20,20,3)) #has to be the same dimension as images fed into the model to predict the label
l=[img_resize.flatten()]
probability=model.predict_proba(l)
for ind,val in enumerate(classes):
    print(f'{val} = {probability[0][ind]*100}%') #predicted probability of classes the image is
print("The predicted image is : "+ classes[model.predict(l)[0]]) #predict the first most likely class of the image
```



Bagel = 8.161424174701217%  
Baguette = 11.48045944535891%  
Brioche = 22.86355032550705%  
Challah = 19.290898862260615%  
Cheese Bread = 5.722914256571526%  
Ciabatta = 4.309797125733436%  
Cinnamon Raisin = 8.09074177568979%  
Focaccia = 5.056570326226335%  
Wheat Grain Bread = 12.337912631605068%  
White Bread = 2.685731076346051%  
The predicted image is : Challah

## Model

BreadNetAlex:

Input data that has been loaded with AlexNet features (256 layers), outputs data as 10 layers corresponding to 10 classes.

#### BreadNetCNN:

Input data that has been normalized and intensity channel removed (3 layers, RGB), outputs data as 10 layers corresponding to 10 classes.

```
In [ ]: import os
import torchvision.models
alexnet = torchvision.models.alexnet(pretrained=True)

data_dir = '/content/gdrive/MyDrive/APS360 Project Group/Data/Alexnet features'
data_directory = '/content/gdrive/MyDrive/APS360 Project Group/Data/Images'
transform = transforms.Compose([transforms.ToTensor(), transforms.Resize((224,224))])
full_dataset = torchvision.datasets.ImageFolder(data_directory, transform=transform)

trainset, val_and_test = torch.utils.data.random_split(full_dataset, [int(0.8 * len(full_dataset)), int(0.2 * len(full_dataset))])
valset, testset = torch.utils.data.random_split(val_and_test, [int(0.5 * len(val_and_test)), int(0.5 * len(val_and_test))])

classes = ('Bagel', 'Baguette', 'Brioche', 'Challah', 'Cheese Bread', 'Ciabatta', 'Cinnamon Raisin', 'Focaccia', 'Wheat Grain Bread', 'White Bread')

from PIL import PngImagePlugin
LARGE_ENOUGH_NUMBER = 100
PngImagePlugin.MAX_TEXT_CHUNK = LARGE_ENOUGH_NUMBER * (1024**2)

n = 0
print('Training Data Converting...')
for i, label in iter(trainset):
    features = alexnet.features(i)
    features_tensor = torch.from_numpy(features.detach().numpy())

    folder_name = data_dir + '/train/' + str(classes[label])
    torch.save(features_tensor.squeeze(0), folder_name + '/' + str(n) + '.tensor')
    n += 1
print('Training Data Completed')

n = 0
print('Validation Data Converting...')
for i, label in iter(valset):
    features = alexnet.features(i)
    features_tensor = torch.from_numpy(features.detach().numpy())

    folder_name = data_dir + '/val/' + str(classes[label])
    torch.save(features_tensor.squeeze(0), folder_name + '/' + str(n) + '.tensor')
    n += 1
print('Validation Data Completed')

n = 0
print('Test Data Converting...')
for i, label in iter(testset):
    features = alexnet.features(i)
    features_tensor = torch.from_numpy(features.detach().numpy())

    folder_name = data_dir + '/test/' + str(classes[label])
    torch.save(features_tensor.squeeze(0), folder_name + '/' + str(n) + '.tensor')
    n += 1
print('Test Data Completed')

/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be
removed in the future, please use 'weights' instead.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or 'None' for 'weights' are
deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing 'weights=AlexNet_Weights.IMAGENET1K_V1'. You can
also use 'weights=AlexNet_Weights.DEFAULT' to get the most up-to-date weights.
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/alexnet-owt-7be5be79.pth" to /root/.cache/torch/hub/checkpoints/alexnet-owt-7be5be79.pth
100%|██████████| 233M/233M [00:04<00:00, 55.5MB/s]
Training Data Converting...

/usr/local/lib/python3.10/dist-packages/torchvision/transforms/functional.py:1603: UserWarning: The default value of the antialias parameter of all th
e resizing transforms (Resize(), RandomResizedCrop(), etc.) will change from None to True in v0.17, in order to be consistent across the PIL and Tens
or backends. To suppress this warning, directly pass antialias=True (recommended, future default), antialias=None (current default, which means False f
or Tensors and True for PIL), or antialias=False (only works on Tensors - PIL will still use antialiasing). This also applies if you are using the inf
erence transforms from the models weights: update the call to weights.transforms(antialias=True).
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/PIL/Image.py:996: UserWarning: Palette images with Transparency expressed in bytes should be converted to RGBA
images
  warnings.warn(
Training Data Completed
Validation Data Converting...
Validation Data Completed
Test Data Converting...
Test Data Completed

In [ ]: #saved pretrained features to separate folders with .tensor extensions
train_features = torchvision.datasets.DatasetFolder(data_dir + '/train/', loader = torch.load, extensions = ('.tensor'))
val_features = torchvision.datasets.DatasetFolder(data_dir + '/val/', loader = torch.load, extensions = ('.tensor'))
test_features = torchvision.datasets.DatasetFolder(data_dir + '/test/', loader = torch.load, extensions = ('.tensor'))

batch_size = 256
num_workers = 0

#loading pretrained data
train_loader_pre = torch.utils.data.DataLoader(train_features, batch_size = batch_size, num_workers = num_workers, shuffle = True)
val_loader_pre = torch.utils.data.DataLoader(val_features, batch_size = batch_size, num_workers = num_workers, shuffle = True)
test_loader_pre = torch.utils.data.DataLoader(test_features, batch_size = batch_size, num_workers = num_workers, shuffle = True)

In [ ]: class BreadNetAlex(nn.Module):
def __init__(self):
super(BreadNetAlex, self).__init__()
self.name = "Alex"
self.conv1 = nn.Conv2d(256, 256, 3, padding=1) #AlexNet 256 input
self.pool = nn.MaxPool2d(2,2)
self.fc1 = nn.Linear(256*3*3, 64)
self.fc2 = nn.Linear(64, 10) #10 classes of bread
```

```

def forward(self, x):
    x = F.relu(self.conv1(x))
    x = self.pool(x)
    x = x.view(-1, 256*3*3)
    x = F.relu(self.fc1(x))
    x = self.fc2(x)
    x.squeeze(1)
    return x

#This model gives the option to not use pre-trained features in training, can compare more hyperparameter modifications with the features in AlexNet
class BreadNetCNN(nn.Module):
    def __init__(self):
        super(BreadNetCNN, self).__init__()
        self.name = "CNN"
        self.conv1 = nn.Conv2d(3, 5, 3, stride=2, padding=1)
        self.pool = nn.MaxPool2d(2,2)
        self.conv2 = nn.Conv2d(5, 5, 7, stride=2, padding=1)
        self.fc1 = nn.Linear(5*13*13, 64)
        self.fc2 = nn.Linear(64, 10)

    def forward(self,x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 5*13*13)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        x.squeeze(1)
        return x

```

In [ ]:

## Possibly Helpful Functions - alice

**get\_accuracy(model, data\_loader):**

Depending on if model is pretrained with AlexNet or CNN, data\_loader needs to be loaded according to the model selected. Returns decimal for accuracy.

**get\_loss(model, data\_loader, criterion):**

data\_loader needs to correspond with model selected. Returns loss as a float.

```

In [ ]: def get_accuracy(model, train_loader, valid_loader, train=False):
    if train:
        data = train_loader
    else:
        data = valid_loader

    correct = 0
    total = 0
    for imgs, labels in data:
        output = model(imgs)

        #select index with maximum prediction score
        pred = output.max(1, keepdim=True)[1]
        correct += pred.eq(labels.view_as(pred)).sum().item()
        total += imgs.shape[0]
    return correct / total

def get_loss(model, data_loader, criterion):
    total_loss = 0.0
    loss = 0.0
    eval_mod = model.eval()

    for i, labels in data_loader:
        output = eval_mod(i)
        loss = criterion(output, labels)
        total_loss += loss.item()

    return (float(total_loss)/len(data_loader))

def get_model_name(name, batch_size, learning_rate, epoch):
    path = 'model_{0}_bs{1}_lr{2}_epoch{3}'.format(name, batch_size, learning_rate, epoch)

    return path

```

In [ ]:

```

In [ ]: #training (name, batch_size, learning_rate, epoch)

#Alice's training code, input is dataset NOT loader
#can input CNN or AlexNet dataset

def train(model, train_dataset, valid_dataset, batch_size=64, learning_rate=0.001, num_epochs=1):
    train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
    valid_loader = torch.utils.data.DataLoader(valid_dataset, batch_size=batch_size, shuffle=True)
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(model.parameters(), lr=learning_rate, momentum=0.9)
    print(len(train_loader))
    print(len(valid_loader))
    iters, losses, train_acc, val_acc = [], [], [], []

    # training
    epoch = 0 # the number of iterations
    for epoch in range(num_epochs):
        for imgs, labels in iter(train_loader):
            out = model(imgs) # forward pass
            loss = criterion(out, labels) # compute the total loss

```

```

        loss.backward()           # backward pass (compute parameter updates)
        optimizer.step()          # make the updates for each parameter
        optimizer.zero_grad()     # a clean up step for PyTorch

    # save the current training information
    iters.append(epoch)
    losses.append(float(loss)/batch_size) # compute *average* loss
    train_acc.append(get_accuracy(model, train_loader, valid_loader, train=True)) # compute training accuracy
    print("epoch number ", epoch+1, "accuracy: ", train_acc[epoch])
    val_acc.append(get_accuracy(model, train_loader, valid_loader, train=False)) # compute validation accuracy
    model_path = "model_{0}_bs{1}_lr{2}_epoch{3}".format(model.name, batch_size, learning_rate, epoch)
    torch.save(model.state_dict(), model_path)

# plotting
plt.title("Training Curve")
plt.plot(iters, losses, label="Train")
plt.xlabel("Iterations")
plt.ylabel("Loss")
plt.show()

plt.title("Training Curve")
plt.plot(iters, train_acc, label="Train")
plt.plot(iters, val_acc, label="Validation")
plt.xlabel("Iterations")
plt.ylabel("Training Accuracy")
plt.legend(loc='best')
plt.show()

print("Final Training Accuracy: {}".format(train_acc[-1]))
print("Final Validation Accuracy: {}".format(val_acc[-1]))

```

## Training for Hyperparameters

adjusting the batch size, learning rate and number of epochs in order to select the best model in terms of accuracy.

```


In [ ]: #training alexnet model
model = BreadNetAlex()


train(model, train_features, val_features, batch_size = 256, learning_rate = 0.01, num_epochs = 6)

```

33  
4

epoch number	accuracy
1	0.6114271879917685
2	0.8365815276600895
3	0.9339063067425252
4	0.9611427187991769
5	0.9920106524633822
6	0.9968526812734536



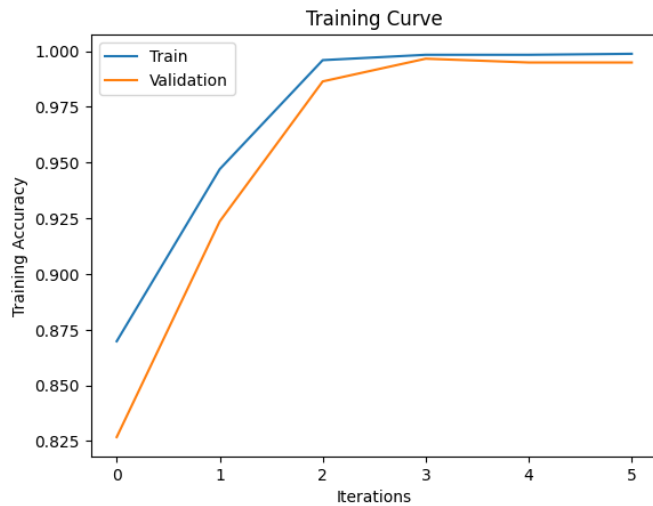



Final Training Accuracy: 0.9968526812734536  
 Final Validation Accuracy: 0.9893992932862191

```
In [ ]: #training alexnet model
model = BreadNetAlex()

train(model, train_features, val_features, batch_size = 64, learning_rate = 0.01, num_epochs = 6)
```



```
105
10
epoch number 1 accuracy: 0.8698189435882089
epoch number 2 accuracy: 0.9470297770462367
epoch number 3 accuracy: 0.9959598982492892
epoch number 4 accuracy: 0.9983540326200808
epoch number 5 accuracy: 0.9983540326200808
epoch number 6 accuracy: 0.9988029328146042
```



Final Training Accuracy: 0.9988029328146042  
 Final Validation Accuracy: 0.9949066213921901

```
In [ ]: model = BreadNetAlex()
train(model, train_features, val_features, batch_size = 128, learning_rate = 0.01, num_epochs = 6)
```

```
53
5
epoch number 1 accuracy: 0.7382911865928475
epoch number 2 accuracy: 0.9001945234176268
epoch number 3 accuracy: 0.9748615891066886
epoch number 4 accuracy: 0.9881789615442167
epoch number 5 accuracy: 0.9979051324255573
epoch number 6 accuracy: 0.9932664970821488
```

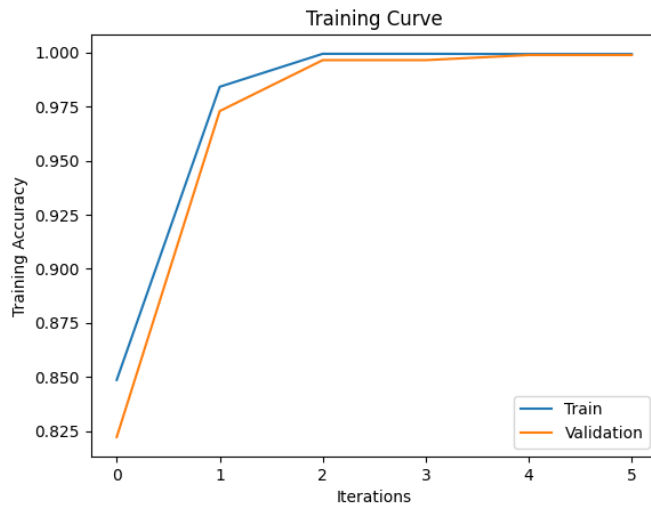



```
Final Training Accuracy: 0.9932664970821488
Final Validation Accuracy: 0.9830220713073005
```

```
In [ ]: model = BreadNetAlex()
train(model, train_features, val_features, batch_size = 32, learning_rate = 0.01, num_epochs = 6)
#final training curve
```

```
259
27
epoch number 1 accuracy: 0.8485655489650163
epoch number 2 accuracy: 0.9841423556470161
epoch number 3 accuracy: 0.999394746398741
epoch number 4 accuracy: 0.999394746398741
epoch number 5 accuracy: 0.9992736956784893
epoch number 6 accuracy: 0.9992736956784893
```







Final Training Accuracy: 0.9992736956784893  
 Final Validation Accuracy: 0.9988221436984688

```
In [ ]: #training for batch_size = 256, learning_rate = 0.01, num_epochs = 6
use_cuda = True
model = BreadNetCNN()
if use_cuda and torch.cuda.is_available():
    model.cuda()
    print('CUDA is available! Training on GPU ...')
else:
    print('CUDA is not available. Training on CPU ...')

#proper model
train(model, train_data, val_data, batch_size = 256, learning_rate = 0.01, num_epochs=6)
```

```
CUDA is not available. Training on CPU ...
6
1
epoch number 1 accuracy: 0.098
epoch number 2 accuracy: 0.11666666666666667
epoch number 3 accuracy: 0.114
epoch number 4 accuracy: 0.114
epoch number 5 accuracy: 0.114
epoch number 6 accuracy: 0.114
```



Final Training Accuracy: 0.114  
Final Validation Accuracy: 0.11413043478260869

```
In [ ]: #training for batch_size = 128, learning_rate = 0.008, num_epochs = 6
use_cuda = True
model = BreadNetCNN()
if use_cuda and torch.cuda.is_available():
    model.cuda()
    print('CUDA is available! Training on GPU ...')
else:
    print('CUDA is not available. Training on CPU ...')

#proper model
train(model, train_data, val_data, batch_size = 128, learning_rate = 0.008, num_epochs=6)
```

```
In [ ]: #training for batch_size = 64, learning_rate = 0.005, num_epochs = 6
use_cuda = True
model = BreadNetCNN()
if use_cuda and torch.cuda.is_available():
    model.cuda()
    print('CUDA is available! Training on GPU ...')
else:
    print('CUDA is not available. Training on CPU ...')

#proper model
train(model, train_data, val_data, batch_size = 64, learning_rate = 0.005, num_epochs=6)
```

```
In [ ]: #training for batch_size = 128, learning_rate = 0.005, num_epochs = 10
use_cuda = True
model = BreadNetCNN()
if use_cuda and torch.cuda.is_available():
    model.cuda()
    print('CUDA is available! Training on GPU ...')
else:
    print('CUDA is not available. Training on CPU ...')

#proper model
train(model, train_data, val_data, batch_size = 128, learning_rate = 0.005, num_epochs=10)
```

```
In [ ]: #training for batch_size = 64, learning_rate = 0.01, num_epochs = 10
use_cuda = True
model = BreadNetCNN()
if use_cuda and torch.cuda.is_available():
    model.cuda()
    print('CUDA is available! Training on GPU ...')
else:
    print('CUDA is not available. Training on CPU ...')
```

```
#proper model
train(model, train_data, val_data, batch_size = 64, learning_rate = 0.01, num_epochs=10)
```

```
In [ ]: #training for batch_size = 256, learning_rate = 0.01, num_epochs = 20
use_cuda = True
model = BreadNetCNN()
if use_cuda and torch.cuda.is_available():
    model.cuda()
    print('CUDA is available! Training on GPU ...')
else:
    print('CUDA is not available. Training on CPU ...')

#proper model
train(model, train_data, val_data, batch_size = 256, learning_rate = 0.01, num_epochs=20)
```

```
In [ ]: #training for batch_size = 32, learning_rate = 0.01, num_epochs = 6
use_cuda = True
model = BreadNetCNN()
if use_cuda and torch.cuda.is_available():
    model.cuda()
    print('CUDA is available! Training on GPU ...')
else:
    print('CUDA is not available. Training on CPU ...')

#proper model
train(model, train_data, val_data, batch_size = 32, learning_rate = 0.01, num_epochs=6)
```

## Getting Qualitative Results

```
In [ ]: from sklearn.metrics import confusion_matrix
import seaborn as sn
import pandas as pd

model = BreadNetAlex()
model_path = "model_{0}_bs{1}_lr{2}_epoch{3}".format("Alex", 32, 0.01, 5)
state = torch.load(model_path)
model.load_state_dict(state)

y_pred = []
y_true = []

train_alex_loader = torch.utils.data.DataLoader(train_features, batch_size=batch_size, num_workers=num_workers, shuffle=True) #or 16


# iterate over train data
for inputs, labels in train_alex_loader:
    output = model(inputs) # Feed Network

    output = (torch.max(torch.exp(output), 1)[1]).data.cpu().numpy()
    y_pred.extend(output) # Save Prediction

    labels = labels.data.cpu().numpy()
    y_true.extend(labels) # Save Truth

# Build confusion matrix
cf_matrix = confusion_matrix(y_true, y_pred)
df_cm = pd.DataFrame(cf_matrix / np.sum(cf_matrix, axis=1)[:], None, index = [i for i in classes],
                    columns = [i for i in classes])
plt.figure(figsize = (12,7))
sn.heatmap(df_cm, annot=True)
```

```
Out[ ]: <Axes: >
```



```
In [ ]: # sample prediction of a random image
from matplotlib.pyplot import imread
import PIL, torch, torchvision

test_image_path = '/content/gdrive/MyDrive/APS360 Project Group/Data/Images/Bagel/plain+bagel.jpg'
img_show=imread(test_image_path)
plt.imshow(img_show)
plt.show()

transform = transforms.Compose([transforms.ToTensor(), transforms.Resize((224,224))])
from PIL import Image
img = Image.open(test_image_path)


img_t = transform(img)

features = alexnet.features(img_t.float())
new_img = torch.from_numpy(features.detach().numpy())
alex_img = new_img.squeeze(0)

out = model(alex_img)
prob = F.softmax(out, dim=1)

pred_label = prob.max(1, keepdim=True)[1].item()
#pred_label = out.max(1, keepdim=True)[1].item()
print('predicted label: ', classes[pred_label], 'with a probability of', prob.max(1, keepdim=True)[0].item()) #change prob-->out

for ind,val in enumerate(classes):
    print(f'{val} = {prob[0][ind]*100}%') #predicted probability of classes the image is
```



```

predicted label: Bagel with a probability of 0.9999986886978149
Bagel = 99.99987030029297%
Baguette = 2.3568477445223834e-06%
Brioche = 0.00012918759603053331%
Challah = 1.196719807694535e-07%
Cheese Bread = 4.420855148623559e-08%
Ciabatta = 8.836963799607744e-11%
Cinnamon Raisin = 1.180573917736183e-06%
Focaccia = 7.26697777508889e-08%
Wheat Grain Bread = 8.082902347439358e-09%
White Bread = 4.173382528449565e-09%

```

## Evaluating the model

```

In [ ]: def get_test_accuracy(model, test_loader):
        data = test_loader

        correct = 0
        total = 0
        for imgs, labels in data:

            output = model(imgs)

            #select index with maximum prediction score
            pred = output.max(1, keepdim=True)[1]
            correct += pred.eq(labels.view_as(pred)).sum().item()
            total += imgs.shape[0]
        return correct / total

In [ ]: best_model = BreadNetAlex()
best_model_path = "model_{0}_bs{1}_lr{2}_epoch{3}".format("Alex", 32, 0.01, 5)
state = torch.load(best_model_path)
best_model.load_state_dict(state)

Out[ ]: <All keys matched successfully>

In [ ]: test_dataloader = torch.utils.data.DataLoader(test_features, batch_size=16, shuffle=True)
test_accuracy = get_test_accuracy(best_model, test_dataloader)
print("test accuracy:", test_accuracy)

test accuracy: 0.996551724137931

```

## Demonstration on new data

```

In [ ]: test_image_path = '/content/gdrive/MyDrive/APS360 Project Group/Data/demo_images/demo_bread(2).jpg'
img_show=imread(test_image_path)
plt.imshow(img_show)
plt.show()

transform = transforms.Compose([transforms.ToTensor(), transforms.Resize((224,224))])
from PIL import Image
img = Image.open(test_image_path)

img_t = transform(img)

features = alexnet.features(img_t.float())
new_img = torch.from_numpy(features.detach().numpy())
alex_img = new_img.squeeze(0)

out = model(alex_img)
prob = F.softmax(out, dim=1)

pred_label = prob.max(1, keepdim=True)[1].item()
print('predicted label: ', classes[pred_label], 'with a probability of', prob.max(1, keepdim=True)[0].item()) #change prob-->out

for ind,val in enumerate(classes):
    print(f'{val} = {prob[0][ind]*100}%') #predicted probability of classes the image is

```



```

predicted label: Wheat Grain Bread with a probability of 0.6434589624404907
Bagel = 0.010713954456150532%
Baguette = 6.0448001022450626e-05%
Brioche = 1.2319226264953613%
Challah = 0.0013820258900523186%
Cheese Bread = 0.23848067224025726%
Ciabatta = 6.584043025970459%
Cinnamon Raisin = 14.003668785095215%
Focaccia = 0.0013894105795770884%
Wheat Grain Bread = 64.34589385986328%
White Bread = 13.582435607910156%

```

```

In [ ]: test_image_path = '/content/gdrive/MyDrive/APS360 Project Group/Data/demo_images/demo_bread(3).jpg'
img_show=imread(test_image_path)
plt.imshow(img_show)
plt.show()

transform = transforms.Compose([transforms.ToTensor(), transforms.Resize((224,224))])
from PIL import Image
img = Image.open(test_image_path)

img_t = transform(img)


features = alexnet.features(img_t.float())
new_img = torch.from_numpy(features.detach().numpy())
alex_img = new_img.squeeze(0)

out = model(alex_img)
prob = F.softmax(out, dim=1)

pred_label = prob.max(1, keepdim=True)[1].item()
print('predicted label: ', classes[pred_label], 'with a probability of', prob.max(1, keepdim=True)[0].item()) #change prob-->out

for ind,val in enumerate(classes):
    print(f'{val} = {prob[0][ind]*100}%') #predicted probability of classes the image is

```



```

predicted label: Wheat Grain Bread with a probability of 0.8391998410224915
Bagel = 6.905612735863542e-06%
Baguette = 1.2455145679268753e-06%
Brioche = 0.04409118741750717%
Challah = 1.8364974039286608e-06%
Cheese Bread = 2.2280519260675646e-05%
Ciabatta = 0.001297139679081738%
Cinnamon Raisin = 0.009951649233698845%
Focaccia = 2.6308391909424245e-08%
Wheat Grain Bread = 83.91998291015625%
White Bread = 16.02464485168457%

```

```

In [2]: %%shell
jupyter nbconvert --to html "/content/Copy_of_Final_Project_Bread.ipynb"

[NbConvertApp] Converting notebook /content/Copy_of_Final_Project_Bread.ipynb to html
[NbConvertApp] Writing 2675631 bytes to /content/Copy_of_Final_Project_Bread.html

```

Out[2]: